

The Mathematics of Diffusion Models: From Score Matching to Stable Diffusion

Orchestration of AI agents

Yanal Serhan & Nell Khoury

Lecturer: Dr. Yoram Segal

June 10, 2026

Abstract

Generated book on The Mathematics of Diffusion Models: From Score Matching to Stable Diffusion.

Contents

Contents	ii
List of Figures	v
List of Tables	vi
1 Introduction to Diffusion Models	1
1.1 Overview of Generative Modeling	1
1.2 Historical Context and Evolution	2
1.3 Mathematical Preliminaries: Stochastic Processes and PDEs	3
1.4 Book Structure and Learning Path	4
2 Fundamentals of Diffusion Processes	6
2.1 Diffusion as a Stochastic Differential Equation (SDE)	6
2.2 Fokker-Planck Equation and Probability Density Evolution .	7
2.3 Ornstein-Uhlenbeck Process and Gaussian Diffusions	8
2.4 Comparative Table: Key Properties of Diffusion Processes .	8
3 Score Matching and its Mathematical Foundations	10
3.1 Definition and Intuition Behind Score Matching	10
3.2 Hyvärinen's Score Matching Estimator	11
3.3 Noise-Conditioned Score Matching	11

3.4	Comparative Table: Score Matching Methods and Their Characteristics	12
4	The Forward and Reverse Diffusion Processes	14
4.1	Forward Diffusion: Noising the Data	14
4.2	Reverse Diffusion: Generative Model as Reverse SDE	15
4.3	Connections to Probability Flow ODEs	15
4.4	Comparative Table: Forward SDE, Reverse SDE, and Probability Flow ODE	16
5	Score-Based Generative Modeling Frameworks	18
5.1	Neural Network Parameterizations of Score Functions	18
5.2	Continuous Noise Schedules and Multi-Scale Training	19
5.3	Practical Sampling Algorithms: DDPM and SDE-Solvers	19
5.4	Comparative Table: Score-Based Model Architectures and Samplers	20
6	The Mathematics Behind Stable Diffusion	22
6.1	Latent Variable Modeling and Dimensionality Reduction	22
6.2	Conditioning Mechanisms and Control	23
6.3	Training Objectives and Stability Guarantees	23
6.4	Comparative Table: Stable Diffusion vs. Classical Diffusion Models	24
7	Theoretical Extensions and Advanced Topics	25
7.1	Likelihood Estimation and ELBO in Diffusion Models	25
7.2	Score-Based Normalizing Flows and Hybrid Models	26
7.3	Diffusion Models for Non-Image Data Domains	26
7.4	Comparative Table: Advanced Diffusion Model Variants	27
8	Practical Implementation and Computational Considerations	28
8.1	Numerical Methods for SDE and ODE Solvers	28
8.2	Efficient Training Techniques and Hardware Optimizations	29

CONTENTS	iv
8.3 Evaluation Metrics and Benchmarking Strategies	29
8.4 Comparative Table: Numerical and Hardware Techniques for Diffusion Models	30
9 Case Studies and Applications of Diffusion Models	31
9.1 Image Synthesis and Editing with Stable Diffusion	31
9.2 Audio Generation and Enhancement via Diffusion	32
9.3 Scientific Data Modeling and Simulation	32
9.4 Comparative Table: Applications Across Domains	32
10 Future Directions and Open Challenges	34
10.1 Theoretical Open Problems in Score Matching and Diffusion	34
10.2 Enhancing Efficiency and Scalability	34
10.3 Ethical Considerations and Responsible AI	35
10.4 Comparative Table: Research Challenges and Potential So- lutions	35
Bibliography	38
11 Hebrew-English BiDi Section	39
11.1 Advanced BiDi Text Integration	39
11.2 Mathematical Formulation in RTL Context	40
System Architecture Diagrams	42
Full Agent Prompts	43
Raw Research Corpus	44
Pipeline Run Statistics	45

List of Figures

4.1	Comparison of Forward SDE, Reverse SDE, and Probability Flow ODE in diffusion models. This figure visually summarizes the mathematical forms, purposes, and qualitative sampling complexities of the three key stochastic/deterministic processes underlying diffusion-based generative modeling.	17
5.1	Noise scheduling schemes employed in score-based diffusion models. The chart illustrates typical continuous noise schedules used to condition neural networks for score matching, including linear, cosine, and sigmoid schedules. Different scheduling shapes impact the training stability and sample quality in diffusion models.	20
5.2	Trade-off between the number of sampling steps and sample quality across popular score-based generative methods. This figure plots hypothetical sample quality against the number of sampling steps required, comparing DDPM, Score SDE, and Probability Flow ODE sampling approaches.	21

List of Tables

2.1	Comparison of classical diffusion processes	9
3.1	Comparison of Score Matching Approaches	12
4.1	Comparison of Forward and Reverse Diffusion Processes	16
5.1	Comparing Score-Based Models and Sampling Methods	21
6.1	Comparative Analysis of Stable Diffusion vs. Classical Diffusion Models	24
7.1	Comparative Overview of Advanced Diffusion Model Variants	27
8.1	Numerical and Hardware Optimization Techniques	30
9.1	Cross-Domain Diffusion Model Applications	33
10.1	Summary of Open Challenges and Research Directions	35

Introduction to Diffusion Models

1.1 Overview of Generative Modeling

Generative modeling refers to a class of methods in machine learning designed to model the underlying probability distribution of observed data. The primary goal is to generate new, plausible samples that resemble the original dataset. Fundamentally, generative models capture complex data distributions to reproduce intricate patterns seen in images, audio, or text. Key categories include explicit density models, such as Variational Autoencoders (VAEs), Autoregressive models (e.g., PixelCNN), and implicit density models like Generative Adversarial Networks (GANs). Each adopts a unique mathematical framework for learning and sampling latent structures.

Diffusion models constitute a relatively recent and promising branch of generative modeling. They work by gradually corrupting data with noise through a forward diffusion process and then learning a reverse denoising process that reconstructs data samples from noise. Unlike GANs, diffusion models avoid adversarial training instabilities. Compared to VAEs, they tend to generate sharper samples due to their explicit denoising approach. These models leverage stochastic differential equations (SDEs) and are mathematically grounded in diffusion processes from physics and probability theory.

The impact of diffusion models is far-reaching. They have demonstrated

state-of-the-art performance in image synthesis, surpassing other generative approaches in terms of sample diversity and quality. Important applications include image editing, super-resolution, molecular generation, and audio synthesis. Their principled probabilistic formulation allows for rigorous theoretical grounding and robust extensions, making them a focal point in current machine learning research **Song2021**.

1.2 Historical Context and Evolution

Diffusion processes initially emerged in the early 20th century within the domain of stochastic analysis. Rooted in Brownian motion studies by Einstein and others, these processes characterized random particle motion in continuous time. The mathematical formalism of diffusion as solutions to stochastic differential equations was later developed by Itô and Stratonovich, providing a foundational toolkit for modern stochastic calculus.

Initially, diffusion was a physical phenomenon describing heat transfer, molecular motion, and other transport phenomena. Its probabilistic interpretation allowed modeling of evolving probability densities instead of deterministic trajectories, encoded by partial differential equations like the Fokker-Planck equation. These insights laid the groundwork for probabilistic modeling frameworks expressing uncertainty and dynamics via diffusion.

The transition to machine learning began with efforts to adapt diffusion processes to represent complex data distributions. Early probabilistic models viewed data as generated through latent stochastic dynamics. Score matching, introduced by Hyvärinen in 2005, was an essential milestone, enabling parameter estimation for unnormalized models by matching gradients of log densities ("scores"). This technique facilitated leveraging diffusion-inspired noise processes for generative modeling without requiring explicit likelihood calculations.

Over the last decade, advances in neural network architectures, increased computational power, and algorithmic innovations have propelled diffusion models from theory to practical tools. Landmark works by Sohl-Dickstein et al. **Sohl-Dickstein2015** and later by Ho et al. **Ho2020** in-

roduced discrete-time diffusion probabilistic models. Subsequent research extended these to continuous-time stochastic differential equations, culminating in scalable, high-fidelity models encompassing the modern class typified by Stable Diffusion **Rombach2022**.

1.3 Mathematical Preliminaries: Stochastic Processes and PDEs

The mathematical core of diffusion models rests on stochastic processes and partial differential equations (PDEs). A stochastic process is a collection of random variables indexed by time, often representing the evolution of a system influenced by randomness. Brownian motion W_t is the canonical continuous-time stochastic process with independent, normally distributed increments and continuous paths. It models pure diffusion without drift, fundamental for constructing more complex dynamics.

Langevin dynamics introduce deterministic drift alongside stochastic noise, described by the stochastic differential equation (SDE):

$$dX_t = b(X_t)dt + \sigma dW_t,$$

where X_t is the state at time t , $b(\cdot)$ is the drift vector field, σ controls diffusion intensity, and W_t denotes Brownian motion. Solutions to SDEs characterize the probabilistic evolution of X_t , capturing both randomness and directed movement.

The associated PDE governing the probability density $p(x, t)$ of the process is the Fokker-Planck equation:

$$\frac{\partial p}{\partial t} = -\nabla \cdot (bp) + \frac{1}{2} \nabla^2 : (Dp),$$

where $D = \sigma\sigma^T$ is the diffusion matrix, $\nabla \cdot$ is divergence, and $\nabla^2 :$ denotes the Laplacian operator applied coordinate-wise. This equation describes how the probability density diffuses and drifts over time.

For this book, the primary assumptions include the time-continuity of stochastic processes, well-behaved drift and diffusion coefficients (Lipschitz

continuity to ensure existence and uniqueness of solutions), and smoothness conditions to justify differentiability of densities. Notation such as $\nabla_x \log p(x)$ will denote the score function—gradient of the log density—central to score matching.

1.4 Book Structure and Learning Path

This book unfolds progressively from fundamental theory to cutting-edge practice in diffusion-based generative modeling. Chapters 1 and 2 introduce diffusion processes mathematically, ensuring a solid foundation in stochastic calculus and PDEs. Chapter 3 delves into score matching, revealing the statistical underpinnings enabling parameter estimation without explicit likelihoods.

The middle chapters (4–6) focus on the forward and reverse diffusion processes alongside the architectures and algorithms that implement score-based models. These bridge theory to applications culminating in detailed analysis of Stable Diffusion—a state-of-the-art model achieving scalable, high-quality generation.

Subsequent chapters explore advanced theoretical topics, practical implementation details, and diverse applications spanning images, audio, and scientific data. The concluding chapter discusses open challenges and future research directions, inviting readers to contribute to evolving frontiers.

The book is designed for graduate students, researchers, and practitioners with a background in probability, calculus, and basic machine learning. Recommended prerequisites include familiarity with stochastic calculus, differential equations, and neural networks. Parallel reading of foundational texts in stochastic processes **Oksendal2003** and generative modeling **Goodfellow2016** will accelerate comprehension.

Chapter 1 Summary: In this chapter, we positioned diffusion models within the landscape of generative modeling, introducing their unique mathematical basis in stochastic processes and PDEs. We traced their evolution from physical diffusion phenomena to machine learning breakthroughs. Lastly, we outlined the book’s roadmap, preparing readers for an in-depth

mathematical exploration of diffusion models and their practical incarnations.

—

Fundamentals of Diffusion Processes

2.1 Diffusion as a Stochastic Differential Equation (SDE)

Diffusion processes are rigorously modeled as stochastic differential equations (SDEs), which describe continuous-time random motions with drift and stochastic components. An SDE formalizes system evolution as:

$$dX_t = b(X_t, t)dt + \sigma(X_t, t)dW_t,$$

where $X_t \in \mathbb{R}^d$ is the state at time t , $b : \mathbb{R}^d \times [0, T] \rightarrow \mathbb{R}^d$ is the drift function representing deterministic influence, $\sigma : \mathbb{R}^d \times [0, T] \rightarrow \mathbb{R}^{d \times m}$ is the diffusion coefficient matrix scaling the noise, and W_t is an m -dimensional Brownian motion.

The drift term b provides a directional bias or deterministic force guiding X_t , shaping the expected trajectory. The diffusion term σdW_t introduces random fluctuations modeling uncertainty or noise, scaled both by σ and the increments of Brownian motion (which are Gaussian).

Forward-time SDEs describe natural stochastic dynamics, such as the process of gradually adding noise to data. Reverse-time SDEs play a vital role in generative modeling, allowing reconstruction by reversing the noise addition process. The reverse SDE for a given forward SDE can be expressed as **Anderson1982; Chen2021**:

$$dX_t = [b(X_t, t) - \sigma(X_t, t)\sigma(X_t, t)^T \nabla_x \log p_t(X_t)] dt + \sigma(X_t, t) d\bar{W}_t,$$

where p_t is the marginal distribution at time t , and $\bar{W}_t$ is a reverse-time Brownian motion. The reverse SDE requires knowledge or estimation of the score function $\nabla_x \log p_t$.

Understanding these forward and reverse SDE formulations is crucial for diffusion models, bridging stochastic process theory with generative sampling mechanics.

2.2 Fokker-Planck Equation and Probability Density Evolution

Every diffusion process modeled by an SDE corresponds to an evolution of its probability density function $p(x, t)$, governed by the Fokker-Planck (FP) equation. The FP equation describes how $p(x, t)$ changes over time due to both drift and diffusion components:

$$\frac{\partial p}{\partial t}(x, t) = - \sum_{i=1}^d \frac{\partial}{\partial x_i} (b_i(x, t)p(x, t)) + \frac{1}{2} \sum_{i,j=1}^d \frac{\partial^2}{\partial x_i \partial x_j} ([D(x, t)]_{ij} p(x, t)),$$

where $D(x, t) = \sigma(x, t)\sigma(x, t)^T$ is the diffusion matrix, typically positive semidefinite.

This PDE generalizes the classical heat equation, which models heat flow as a pure diffusion process with zero drift and constant diffusion coefficient. Thus, the FP equation serves as a fundamental link between microscopic stochastic dynamics and macroscopic density evolution.

Stationary distributions correspond to steady-state solutions $p_\infty(x)$ satisfying

$$0 = -\nabla \cdot (b(x)p_\infty(x)) + \frac{1}{2} \nabla^2 : (D(x)p_\infty(x)),$$

representing equilibrium densities towards which the process converges.

Solving the FP equation explicitly is challenging except for simple cases. Nevertheless, it provides theoretical insight and underpins score-based modeling since the score function involves gradients of log densities evolved under the FP dynamics.

2.3 Ornstein-Uhlenbeck Process and Gaussian Diffusions

The Ornstein-Uhlenbeck (OU) process serves as a fundamental example of a linear diffusion process with analytical tractability. It solves the SDE:

$$dX_t = -\theta X_t dt + \sigma dW_t,$$

where $\theta > 0$ is the mean reversion rate, pulling the process back toward zero, and σ is the diffusion strength.

The OU process admits a Gaussian stationary distribution $\mathcal{N}(0, \frac{\sigma^2}{2\theta}I)$, demonstrating stationarity and ergodicity properties. Its transition density is explicitly known as a Gaussian with mean $e^{-\theta t}x_0$ and covariance matrix:

$$\Sigma_t = \frac{\sigma^2}{2\theta}(1 - e^{-2\theta t})I.$$

This process is a cornerstone for constructing and understanding more complex diffusion models because it blends deterministic damping with stochastic noise in a mathematically manageable way.

General Itô diffusions extend this framework by allowing nonlinear and time-dependent drift $b(x, t)$ and diffusion $\sigma(x, t)$ terms, but typically lack closed-form solutions, requiring numerical approaches or approximations.

2.4 Comparative Table: Key Properties of Diffusion Processes

Brownian motion exemplifies pure diffusive motion without drift, resulting in ever-spreading densities without equilibrium. The OU process adds

Process	Drift	Diffusion	Stationary Dist.	Analytic Solution
Brownian Motion	0	σ (constant)	None (non-stationary)	Yes
Ornstein-Uhlenbeck	$-\theta x$	σ (constant)	Gaussian	Yes
General Itô Diffusion	$b(x, t)$	$\sigma(x, t)$	Possibly	Generally No

Table 2.1: Comparison of classical diffusion processes

restoring drift allowing a stationary Gaussian distribution. General Itô diffusions encompass a broad class, including nonlinear dynamics and state-dependent noise, complicating both analysis and simulation.

Understanding these classes aids in comprehending how diffusion models capture complex data distributions by combining drift and diffusion appropriately, laying the foundation for score-based likelihood-free generative modeling.

Chapter 2 Summary: Chapter 2 elaborated the mathematical backbone of diffusion models using stochastic differential equations and their corresponding PDEs. By examining Brownian motion and Ornstein-Uhlenbeck processes, it illustrated the principles of stationarity, ergodicity, and analytic tractability. This understanding is pivotal for grasping subsequent chapters that leverage stochastic calculus for generative modeling.

Score Matching and its Mathematical Foundations

3.1 Definition and Intuition Behind Score Matching

Score matching is a statistical technique designed to estimate parameters of unnormalized probability density models by matching their gradients of the log density, known as score functions. The score function for a density $p(x)$ is:

$$s(x) = \nabla_x \log p(x).$$

Unlike likelihood-based methods, score matching bypasses the need to compute normalization constants, often intractable in complex models. The intuition lies in directly fitting the gradient fields of the model's log-density to that of the true data distribution, avoiding explicit density evaluation.

Matching scores aligns the geometry of the model's density contours with the data distribution. This approach connects closely to Fisher divergence, which measures the discrepancy between score functions. Formally, the Fisher divergence between two densities p and q is:

$$D_F(p||q) = \frac{1}{2} \int p(x) \|\nabla_x \log p(x) - \nabla_x \log q(x)\|^2 dx.$$

Minimizing this divergence encourages the model to replicate the shape of the data density without requiring explicit normalization.

Score matching thus provides a promising framework for parameter estimation, particularly for energy-based models where partition functions are difficult to compute. It also serves as the mathematical underpinning for training score-based diffusion models, where the score function guides the generative reverse process.

3.2 Hyvärinen’s Score Matching Estimator

Introduced by Hyvärinen (2005), the original score matching estimator formulates a loss function that depends only on the model’s score function and its derivatives. The score matching loss for model density $p_\theta(x)$ is:

$$J(\theta) = \mathbb{E}_{p_{data}} \left[\frac{1}{2} \|\nabla_x \log p_\theta(x)\|^2 + \Delta_x \log p_\theta(x) \right],$$

where Δ_x denotes the Laplacian operator applied to $\log p_\theta(x)$.

This loss is unbiased w.r.t. the parameter θ and does not require knowledge of the normalization constant. The estimator is consistent and, under regularity conditions, asymptotically efficient.

However, practical application in high-dimensional spaces faces challenges. Computing the score function’s gradients and Laplacians can be expensive or numerically unstable. Additionally, direct use of the original form assumes access to clean data without noise, which can be restrictive. These challenges inspired extensions such as denoising and noise-conditioned variants to handle more realistic settings.

3.3 Noise-Conditioned Score Matching

Noise-conditioned score matching extends the original framework to settings where data are corrupted by varying levels of noise. Instead of estimating the score of clean data distribution $p_{data}(x)$ alone, the model learns the scores of noisy distributions $p_{data_t}(x)$ for noise scales indexed by t .

This approach models conditional score functions:

$$s_\theta(x, t) \approx \nabla_x \log p_{data_t}(x),$$

where noise levels vary continuously. This is essential for training diffusion models, which perturb data with a continuous diffusion process.

Noise conditioning improves robustness and sample quality because the model learns to denoise inputs corrupted at multiple noise intensities. Practically, training employs denoising score matching objectives where noisy samples $\tilde{x}$ are generated, and the network predicts the gradient of the log probability of noisy data given the clean sample. This formulation unifies score matching with denoising autoencoders.

The continuous noise schedules underpin modern diffusion models' success by enabling smooth interpolation across noise levels, facilitating efficient and stable training procedures.

3.4 Comparative Table: Score Matching Methods and Their Characteristics

Method	Loss Objective	Noise Handling	Computational Complexity
Original Score Matching	$\mathbb{E}[\nabla \log p_\theta(x) ^2 + 2\Delta \log p_\theta(x)]$	No noise	Moderate
Denoising Score Matching	$\mathbb{E}[s_\theta(\tilde{x}) - \nabla \log p(\tilde{x} x) ^2]$	Explicit noise	High (due to noise marginalization)
Noise-Conditioned Score Matching	Conditional score approximation for varying noise scales	Explicit, continuous noise schedule	Higher, but more robust

Table 3.1: Comparison of Score Matching Approaches

Original score matching is conceptually straightforward but limited by noise-free assumptions. Denoising score matching improves empirical performance but incurs higher computational costs due to noise marginalization. Noise-conditioned score matching, while the most computationally intensive, offers greater robustness and scalability, instrumental to the diffusion modeling revolution.

Chapter 3 Summary: Chapter 3 introduced the score matching framework as a likelihood-free parameter estimation method. It covered Hyvärinen's original formulation and its evolution to noise-conditioned variants

essential for training modern diffusion models. Through this mathematical foundation, readers are prepared for generative modeling techniques exploiting score functions.

—

The Forward and Reverse Diffusion Processes

4.1 Forward Diffusion: Noising the Data

The forward diffusion process gradually adds noise to the data, transforming it into a simple known distribution, often Gaussian. Formally, starting from data $x_0 \sim p_0(x)$, the forward SDE is:

$$dx = f(x, t)dt + g(t)dw_t,$$

where f represents the drift, g the diffusion coefficient, and w_t is Brownian motion. This process is designed so that as $t \rightarrow T$, the noisy data distribution $p_t(x)$ approaches a tractable prior, typically a standard Gaussian.

This forward diffusion defines a semigroup P_t acting on densities evolving smoothly, with the marginal distributions characterized by the Fokker-Planck equation. The semigroup property ensures that the composition of noise steps corresponds to the integrated dynamics of the SDE.

The forward process is fixed and does not require learning. Its mathematical construction ensures invertibility, laying the groundwork for the reverse process that recovers data from noise.

4.2 Reverse Diffusion: Generative Model as Reverse SDE

The generative capability hinges on running a reverse-time SDE that inverts the forward noising process. The reverse SDE is:

$$dx = [f(x, t) - g(t)^2 \nabla_x \log p_t(x)] dt + g(t) d\bar{w}_t,$$

where $\nabla_x \log p_t(x)$ is the score function of the forward process marginals at time t , and $\bar{w}_t$ is Brownian motion in reverse time.

Since p_t is unknown, the score function is approximated via neural networks trained on score matching objectives. Sampling from this reverse SDE produces synthetic data resembling the original distribution.

Practically, numerical schemes such as Euler-Maruyama approximate these continuous dynamics. This involves discretizing the interval $[0, T]$ and iteratively applying stochastic updates conditioned on score estimates.

The reverse SDE's form elucidates how generative modeling is fundamentally a stochastic denoising process guided by learned gradients, bridging physics-inspired diffusion with neural implicit modeling.

4.3 Connections to Probability Flow ODEs

The probability flow ODE represents a deterministic counterpart to the reverse SDE that evolves the data distribution identically in law. It is defined as:

$$dx = \left[f(x, t) - \frac{1}{2} g(t)^2 \nabla_x \log p_t(x) \right] dt.$$

Unlike the stochastic reverse SDE, the ODE has no noise term, making sampling deterministic.

This equivalence follows from the theory of stochastic flows and probability conservation. Importantly, the probability flow ODE allows computing exact likelihoods using change-of-variable formulas, desirable for model evaluation.

Though deterministic, sampling with ODE solvers can be slower but typically yields higher sample quality and stable trajectories. This connection links score-based generative modeling with normalizing flows and optimal transport theory.

4.4 Comparative Table: Forward SDE, Reverse SDE, and Probability Flow ODE

Process	Mathematical Form	Purpose	Sampling Complexity
Forward SDE	$dx = f(x, t)dt + g(t)dw_t$	Data noising	Simple (simulation)
Reverse SDE	$dx = [f(x, t) - g(t)^2 \nabla \log p_t(x)] dt + g(t)d\bar{w}_t$	Data generation	Moderate to high
Probability Flow ODE	$dx = [f(x, t) - \frac{1}{2}g(t)^2 \nabla \log p_t(x)] dt$	Deterministic sampling	Typically lower

Table 4.1: Comparison of Forward and Reverse Diffusion Processes

The forward SDE is conceptually straightforward and numerically simple but used only in the noising direction. The reverse SDE incorporates score estimates with stochasticity, enabling generative sampling. The probability flow ODE provides a noise-free alternative with favorable theoretical properties but may require more sophisticated solvers.

Chapter 4 Summary: This chapter described diffusion models' central bidirectional mechanism via forward and reverse SDEs and the alternative probability flow ODE. It illuminated the mathematical formulation enabling data generation from noise using learned score functions and contrasted stochastic and deterministic sampling strategies.




Figure 4.1: Comparison of Forward SDE, Reverse SDE, and Probability Flow ODE in diffusion models. This figure visually summarizes the mathematical forms, purposes, and qualitative sampling complexities of the three key stochastic/deterministic processes underlying diffusion-based generative modeling.

Score-Based Generative Modeling Frameworks

5.1 Neural Network Parameterizations of Score Functions

The score function $s_\theta(x, t)$, approximating $\nabla_x \log p_t(x)$, is parametrized by neural networks trained to predict scores at different noise levels t . Architectures often employ denoising score matching networks with deep convolutional U-Net structures focused on images.

The network takes as input the noisy data sample x and the noise level t , encoded via sinusoidal embeddings or learned embeddings to condition the model on continuous time. This conditioning allows the network to flexibly adapt to varying noise scales, crucial for accurate score estimation across the diffusion trajectory.

Training optimizes loss functions aligned with noise-conditioned score matching, minimizing the discrepancy between predicted scores and empirical gradients of the noisy data distributions. Techniques like batch normalization, residual connections, and attention layers improve representational power and training stability.

Advanced designs handle multimodal outputs, scalable architectures, and efficient parameter sharing. The versatility of neural parameterizations underpins the widespread adoption and success of score-based diffu-

sion models.

5.2 Continuous Noise Schedules and Multi-Scale Training

Noise schedules define how noise intensity varies over diffusion time t . Continuous-time noise schedules provide more flexible learning than discrete steps, enabling smooth interpolation and stable gradients.

Common schedules include linear, cosine, and sigmoid shapes, each affecting how the model perceives noise levels across training. Continuous conditioning via time embeddings ensures the network learns scores for an entire spectrum of noise scales rather than discrete levels, enhancing generalization.

Multi-scale training arises naturally by sampling noise times uniformly or with importance weighting, ensuring balanced learning across all noise intensities. This mitigates model bias toward easier noise levels and improves generation quality.

Careful noise schedule design and multi-scale training yield improved sample fidelity, convergence speed, and robustness to noise perturbations, fundamental for practical diffusion modeling.

5.3 Practical Sampling Algorithms: DDPM and SDE-Solvers

Denoising Diffusion Probabilistic Models (DDPM) pioneered discrete-time diffusion sampling by applying a Markov chain of learned denoising steps:

$$x_{t-1} = \frac{1}{\sqrt{1 - \beta_t}} \left(x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z,$$

where β_t controls noise magnitude and ϵ_θ predicts noise components.




Figure 5.1: Noise scheduling schemes employed in score-based diffusion models. The chart illustrates typical continuous noise schedules used to condition neural networks for score matching, including linear, cosine, and sigmoid schedules. Different scheduling shapes impact the training stability and sample quality in diffusion models.

Continuous-time score-based models apply Euler-Maruyama or Predictor-Corrector solvers for reverse SDEs, iteratively refining samples by simulating stochastic dynamics guided by score estimates.

Euler-Maruyama is simple and widely used but may require many steps to ensure accuracy. Predictor-Corrector and higher-order schemes improve sampling quality and stability, though at increased computational cost.

There is a trade-off between speed (fewer steps) and fidelity (more steps). Subsequent research on accelerated samplers aims to reduce sampling burden while maintaining sample quality.

5.4 Comparative Table: Score-Based Model Architectures and Samplers

Chapter 5 Summary: This chapter connected theoretical constructs of score functions with neural network parameterizations and training. It

Model/Sampler	Parameterization	Sampling Method	Sampling Steps
DDPM	U-Net conditioned on noise level	Markov chain of de-noising steps	1000+
Score SDE	Noise-conditional score network	Euler-Maruyama reverse SDE integration	100-1000
Probability Flow ODE	Same as Score SDE	ODE solvers (e.g., Runge-Kutta)	Typically fewer & deterministic

Table 5.1: Comparing Score-Based Models and Sampling Methods




Figure 5.2: Trade-off between the number of sampling steps and sample quality across popular score-based generative methods. This figure plots hypothetical sample quality against the number of sampling steps required, comparing DDPM, Score SDE, and Probability Flow ODE sampling approaches.

discussed noise scheduling, conditioning, and practical sampling algorithms crucial for realizing score-based generative modeling in computational frameworks.

The Mathematics Behind Stable Diffusion

6.1 Latent Variable Modeling and Dimensionality Reduction

Stable Diffusion introduces a key innovation by performing diffusion in latent space rather than pixel space. The data x is encoded into a lower-dimensional latent $z = E(x)$ via a Variational Autoencoder (VAE), compressing and representing data efficiently.

The VAE comprises an encoder network producing a distribution $q_\phi(z|x)$ and a decoder reconstructing x from z . This latent space exhibits lower dimensionality and smoother structure, significantly reducing computational costs of diffusion.

The diffusion process is defined on z , with forward and reverse SDEs applied in latent space. Formally, if $p_0(z)$ denotes the latent data distribution, diffusion transforms it into a simple prior (e.g., standard Gaussian). Sampling is performed reversely in latent space, then decoded to data space.

This decoupling enables more scalable and efficient generative modeling suitable for high-resolution images without prohibitive computational demands.

6.2 Conditioning Mechanisms and Control

A critical feature of Stable Diffusion is conditional generation, especially text-to-image synthesis. Conditioning is realized by embedding textual prompts using pretrained language-image models such as CLIP, which produce dense vector representations c .

Conditioning is incorporated via cross-attention layers inside the U-Net score network acting on latent variables:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V,$$

where the query Q derives from latent features and keys K and values V from conditioning embeddings.

Mathematically, conditioning imposes constraints on the reverse SDE by modulating the score function:

$$s_\theta(z, t, c) \approx \nabla_z \log p_t(z|c),$$

guiding generative trajectories toward samples aligning with instructed semantics.

This formulation allows compositional and flexible control, enabling diverse applications such as image editing and style transfer.

6.3 Training Objectives and Stability Guarantees

Stable Diffusion jointly optimizes reconstruction losses from the VAE and denoising score matching losses for the diffusion model:

$$\mathcal{L} = \mathbb{E}_{x,z,t} \left[\|D(E(x)) - x\|^2 + \lambda_t \|s_\theta(z_t, t, c) - \nabla_z \log p_t(z_t|z_0)\|^2 \right].$$

Regularization terms, such as spectral normalization and gradient clipping, ensure stability during training, mitigating mode collapse and divergence issues.

Theoretical stability arises from balancing latent space smoothness via the VAE and robust score estimation conditioned on noise scales and prompts. Optimization techniques such as adaptive learning rates and noise scheduling contribute to convergence guarantees.

Together, these mathematical ingredients underpin Stable Diffusion’s robust performance, enabling state-of-the-art generation quality with reduced computational demands.

6.4 Comparative Table: Stable Diffusion vs. Classical Diffusion Models

Feature	Classical Diffusion	Stable Diffusion	Computational Cost	Generation Quality
Latent Representation	No	Yes (VAE latent space)	Lower	Higher
Conditioning	Limited	Text embeddings + cross-attention	Moderate	Higher control
Sampling Steps	High (1000+)	Fewer (50-100)	Lower	Comparable or better
Training Data Scale	Moderate	Very large scale datasets	Higher	Improved generalization

Table 6.1: Comparative Analysis of Stable Diffusion vs. Classical Diffusion Models

Stable Diffusion’s latent approach reduces memory and computation while preserving detail. Conditioning mechanisms enable nuanced control absent in classical unconditional diffusion. Efficient sampling further accelerates inference. Leveraging large-scale training datasets enhances generalization and quality.

Chapter 6 Summary: Chapter 6 dissected the mathematical factors behind Stable Diffusion’s effectiveness, emphasizing latent space diffusion, text conditioning via cross-attention, and stable training objectives. These innovations mark a paradigm shift moving beyond classical diffusion models toward scalable, controllable generation.

—

Theoretical Extensions and Advanced Topics

7.1 Likelihood Estimation and ELBO in Diffusion Models

Evaluating the likelihood of data under diffusion models is challenging due to implicit constructions and intractable densities. Variational inference techniques derive an Evidence Lower Bound (ELBO) to approximate likelihoods.

The ELBO is:

$$\log p_0(x_0) \geq \mathbb{E}_q [\log p(x_{0:T}) - \log q(x_{1:T}|x_0)],$$

where q is the forward diffusion (noising) process, and p the learned reverse model. Decompositions yield tractable terms relating to denoising score matching losses and entropy integrals.

Optimizing the ELBO aligns model and data distributions, providing interpretability and a path toward likelihood-based evaluation.

Recent work refines ELBO tightness using continuous-time formulations and importance weighting, enhancing training fidelity and interpretability.

7.2 Score-Based Normalizing Flows and Hybrid Models

Normalizing flows offer exact likelihood evaluation by constructing invertible transformations with tractable Jacobians. Combining flows with score-based models yields hybrid architectures benefiting from both score approximation and exact likelihood.

Mathematically, diffusion can be viewed as a flow in probability space. Score-based flows parameterize transformations guided by score functions within flow frameworks, enabling flexible density modeling with exact likelihood calculation.

These hybrids enhance expressiveness and sample quality, marrying the strengths of both paradigms. Theoretical advances provide proofs for invertibility, stability, and convergence guaranteeing model soundness.

7.3 Diffusion Models for Non-Image Data Domains

Diffusion modeling extends beyond images into audio, graphs, and tabular data, each posing unique challenges.

Audio diffusion must model temporal dependencies and signals at high sampling rates. Conditioning on speech content requires architectures capable of capturing intricate frequency patterns.

Graphs involve structured data with variable topology; adapting noise schedules for discrete, combinatorial structures necessitates specialized diffusion operators and score functions.

Tabular data challenges include heterogeneous feature types and missing values, prompting tailored noise injection and conditioning mechanisms.

Theoretical research addresses domain-specific noise processes and model architectures, broadening diffusion models' applicability.

7.4 Comparative Table: Advanced Diffusion Model Variants

Variant	Key Mathematical Feature	Domain	Advantages
Standard Score-Based	Score function approximation	Images, general	Flexible, powerful
Score-Flow Hybrid	Exact likelihood via flows	Images, tabular	Likelihood evaluation
Continuous-Time Flows	Diffusion as flow ODEs	Continuous data	Deterministic sampling
Domain-Specific Diffusion	Tailored noise processes	Audio, graphs	Application-specific gains

Table 7.1: Comparative Overview of Advanced Diffusion Model Variants

Each variant builds on foundational diffusion concepts, adapting or enhancing mathematical structures to suit specific modeling demands and improve performance or applicability.

Chapter 7 Summary: This chapter surveyed theoretical and applied extensions of diffusion models, including likelihood estimation frameworks, hybrid architectures combining flows and score models, and adaptation to diverse data domains. It highlighted ongoing innovations broadening diffusion modeling horizons.

—

Practical Implementation and Computational Considerations

8.1 Numerical Methods for SDE and ODE Solvers

Efficient and stable solution of forward and reverse processes requires numerical solvers for SDEs and ODEs. The Euler-Maruyama method discretizes SDEs with steps:

$$x_{k+1} = x_k + f(x_k, t_k)\Delta t + g(t_k)\Delta W_k,$$

where $\Delta W_k \sim \mathcal{N}(0, \Delta t)$.

While easy to implement, Euler-Maruyama has limited strong order accuracy (0.5) and may require very small step sizes for stability.

Higher-order solvers, such as Runge-Kutta for ODEs and Predictor-Corrector schemes for SDEs, improve accuracy and convergence rates, allowing larger step sizes and faster sampling.

Adaptive step-size methods dynamically vary Δt to balance speed and accuracy but introduce complexity in implementation.

An understanding of stability and convergence for these solvers is essential for reliable model training and sampling.

8.2 Efficient Training Techniques and Hardware Optimizations

Training large diffusion models demands significant computation. Techniques reducing memory and time include mixed precision training, leveraging 16-bit floating point to accelerate matrix operations while preserving accuracy via scaling.

Gradient checkpointing stores intermediate results selectively, trading computation for reduced memory footprint.

Distributed training splits data and model across multiple GPUs or nodes, parallelizing gradient computation and communication to reduce wall time.

Preprocessing data with normalization, augmentation, and careful batching improves convergence and numerical stability.

Hardware-aware optimization, including use of tensor cores and efficient kernels, further accelerates training pipelines.

8.3 Evaluation Metrics and Benchmarking Strategies

Evaluating diffusion models employs metrics quantifying sample quality and diversity. The Fréchet Inception Distance (FID) measures the Wasserstein-2 distance between embeddings of generated and real images, reflecting perceptual similarity.

Inception Score (IS) assesses confidence and variety in generated samples but less correlates with human judgment than FID.

Audio models use Signal-to-Noise Ratio (SNR) and Perceptual Evaluation of Speech Quality (PESQ).

Benchmarking involves standardized datasets, such as CIFAR-10 or ImageNet, with fixed protocols to ensure comparability.

Trade-offs exist between computational cost and metric sensitivity; careful design of evaluation enables informed model development.

8.4 Comparative Table: Numerical and Hardware Techniques for Diffusion Models

Technique	Purpose	Pros	Cons
Euler-Maruyama	Simple SDE solver	Easy to implement	Lower accuracy
Runge-Kutta (higher order)	Improved ODE solver	Better accuracy	Increased computation
Mixed Precision Training	Faster training	Lower memory use	Possible numerical instability
Distributed Training	Scale to large datasets	Shorter wall time	Complex setup

Table 8.1: Numerical and Hardware Optimization Techniques

Choosing appropriate solvers and hardware optimizations depends on trade-offs among accuracy, speed, and resource availability, crucial for practical diffusion model deployment.

Chapter 8 Summary: Chapter 8 offered practical guidelines for implementing diffusion models efficiently, focusing on numerical solvers, training optimizations, and rigorous evaluation metrics. These insights empower practitioners to realize theoretical models robustly and at scale.

—

Case Studies and Applications of Diffusion Models

9.1 Image Synthesis and Editing with Stable Diffusion

Stable Diffusion enables creating photorealistic images from textual prompts through a pipeline: a text prompt is encoded into embeddings, which condition the latent diffusion model. The forward process maps images into latent space via a VAE; the reverse diffusion samples latent variables conditioned on embeddings; finally, decoding produces images.

Latent space manipulations enable image editing and interpolation. For example, linear interpolations between latent vectors correspond to smooth morphing between images—a property stemming from latent space continuity.

Mathematical interpretations characterize how cross-attention guides sample trajectories to relevant semantic modes, affording controllable generation.

Empirical results on benchmarks reveal high-fidelity samples with coherent structure and diverse content, exceeding prior methods in quality-speed trade-offs.

9.2 Audio Generation and Enhancement via Diffusion

Diffusion models applied to audio model waveform probabilities directly or spectral representations. Forward processes inject noise in time or frequency domains, while reverse diffusion denoises toward clean signals.

Conditional diffusion allows speech enhancement by conditioning on noisy audio, learning to map corrupted signals to clean speech distributions.

Challenges arise from non-stationarity and temporal correlations in audio. Architectures incorporate recurrent or convolutional temporal layers.

Mathematical techniques include designing appropriate noise schedules respecting audio dynamics and loss functions tailored for perceptual quality metrics.

9.3 Scientific Data Modeling and Simulation

In physics and chemistry, diffusion models simulate molecular trajectories and configurations, capturing distributions over complex state spaces.

Physically-informed noise processes embed known dynamics, ensuring generated samples obey conservation laws or energy constraints.

Applications include drug discovery, material science, and climate modeling.

Metrics evaluate reconstruction errors, physical validity measures, and uncertainty quantification to assess model reliability.

9.4 Comparative Table: Applications Across Domains

Each domain leverages domain-specific modeling choices and evaluation frameworks, demonstrating diffusion models' adaptability and potency.

Domain	Data Type	Diffusion Model Approach	Evaluation Metrics
Image Generation	Pixel data	Latent diffusion (Stable Diffusion)	FID, IS
Audio Synthesis	Time-series waveform	Conditional score matching	SNR, PESQ
Scientific Simulation	Molecular coordinates	Physics-informed diffusion	RMSE, Physical validity

Table 9.1: Cross-Domain Diffusion Model Applications

Chapter 9 Summary: This chapter connected theory to real-world use cases, showcasing diffusion models’ versatility in image synthesis, audio processing, and scientific simulations. Mathematical rigor informs algorithm design, enabling state-of-the-art applications.

Future Directions and Open Challenges

10.1 Theoretical Open Problems in Score Matching and Diffusion

Despite successes, diffusion models face theoretical challenges. Quantifying approximation errors in score estimation and bounding convergence rates remain open.

Understanding the geometry of loss landscapes could yield improved optimization strategies.

Robustness to noisy or out-of-distribution data requires stronger stability results.

Mathematically rigorous stability and robustness theories would underpin safer and more reliable model deployment.

10.2 Enhancing Efficiency and Scalability

Sampling speed remains a bottleneck, with high step counts limiting real-time applications.

Research explores learned samplers, leveraging neural networks to approximate reverse dynamics more efficiently.

Self-distillation and progressive training paradigms aim to reduce training time.

Integrating diffusion models with other generative frameworks, like GANs or VAEs, may combine complementary strengths.

10.3 Ethical Considerations and Responsible AI

Datasets encode biases propagated into diffusion models, raising fairness concerns.

Enormous computational requirements imply significant environmental impacts, necessitating efficient architectures.

Transparency is limited by model complexity, challenging interpretability and accountability.

Responsible diffusion modeling demands developing tools for bias mitigation, model auditing, and reducing carbon footprints.

10.4 Comparative Table: Research Challenges and Potential Solutions

Challenge	Mathematical/Technical Aspect	Proposed Solutions
Sampling Speed	High number of diffusion steps	Learned samplers, ODE-based methods
Model Robustness	Sensitivity to noise and training data	Regularization, adversarial training
Interpretability	Lack of theoretical transparency	Score function analysis, simpler architectures
Environmental Impact	High computational cost	Efficient architectures, dataset distillation

Table 10.1: Summary of Open Challenges and Research Directions

Addressing these challenges demands combined mathematical innovation, system-level optimization, and ethical consideration.

Chapter 10 Summary: The final chapter outlined key open questions in diffusion research, spanning theory, efficiency, and ethics. It emphasizes a holistic approach marrying mathematical rigor and societal responsibility to steer future development.

References

Bibliography

- [1] Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, M., Ermon, S., & Poole, B. (2021). Score-Based Generative Modeling through Stochastic Differential Equations. *International Conference on Learning Representations*.
- [2] Sohl-Dickstein, J., Weiss, E., Maheswaranathan, N., & Ganguli, S. (2015). Deep Unsupervised Learning using Nonequilibrium Thermodynamics. *Proceedings of the 32nd International Conference on Machine Learning*.
- [3] Ho, J., Jain, A., & Abbeel, P. (2020). Denoising Diffusion Probabilistic Models. *Advances in Neural Information Processing Systems*.
- [4] Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). High-Resolution Image Synthesis with Latent Diffusion Models. *Proceedings of CVPR*.
- [5] Chen, Y., Song, Y., & Ermon, S. (2021). Likelihood-based Generative Modeling with Score-Based Diffusion Models. *Advances in Neural Information Processing Systems*.
- [6] Øksendal, B. (2003). *Stochastic Differential Equations: An Introduction with Applications*. Springer.
- [7] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.

Hebrew-English BiDi Section

This chapter demonstrates advanced mixed Right-To-Left (Hebrew) and Left-To-Right (English) text handling, as well as complex mathematical formatting. Modern LaTeX makes it seamless to switch between left-to-right (LTR) and right-to-left (RTL) paradigms within the same document, or even the same paragraph, which is essential for globalized academic research.

11.1 Advanced BiDi Text Integration

בעידן המודרני של בינה מלאכותית (Artificial Intelligence), אנו עדים להתפתחות מואצת של מודלים שפתיים גדולים (Large Language Models - LLMs). מודלים אלו מאפשרות עיבוד שפה טבעית (Natural Language Processing) ברמה שלא נראתה כמותה בעבר. חשוב לציין כי שילוב של טקסטים אנגליים ועבריים באותו משפט (Inline BiDi) הוא אתגר טכני לא קטן. למשל, כאשר מדברים על Frameworks כגון CrewAI או LangChain, יש לוודא שהכיווניות (Directionality) נשמרת בצורה נאותה. אלגוריתמים מתקדמים לומדים את המבנה התחבירי של מספר שפות במקביל (Multilingual Learning), ומסוגלים לייצר טקסט קוהרנטי וזורם תוך שמירה על כללי הטיפוגרפיה. בנוסף, מערכות אלו נדרשות להתמודד עם אלמנטים נוספים כמו מספרים (12345), משוואות מתמטיות, וציטוטים, כל זאת מבלי לפגוע בקריאות של המסמך.

11.2 Mathematical Formulation in RTL Context

When formulating algorithms for BiDi processing, standard LTR math blocks are often maintained even within RTL paragraphs. The following is a professionally typeset mathematical equation demonstrating a complex cost function for a multilingual model:

$$\mathcal{L}(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^V y_{ij} \log(\hat{y}_{ij}(\theta)) + \lambda \|\theta\|_2^2 \quad (11.1)$$

Where $\mathcal{L}(\theta)$ is the loss, N is the number of samples, and V is the vocabulary size. The inclusion of the regularization term $\lambda \|\theta\|_2^2$ ensures the model does not overfit to specific language syntactic structures.

System Architecture Diagrams

This appendix contains the system architecture diagrams for the CrewAI pipeline.

Full Agent Prompts

This appendix contains the Role, Goal, and Backstory for each agent in the system.

Raw Research Corpus

This appendix contains the annotated source list from the research phase.

Pipeline Run Statistics

This appendix contains the statistics for the pipeline run, including tokens, cost, and time.